

Intro to Fuzzing

How to bug hunting
for fun and profit


- 1 What and Why Fuzzing
- 2 The Zen of Fuzzing
- 3 Fuzzing Kata
- 4 Embedded Fuzzing

What and Why Fuzzing



```
#define Fuzzing
```

“Fuzzing is an automated software testing technique that attempts to find hackable software bugs by randomly feeding invalid and unexpected inputs and data into a program in order to find coding errors and security loopholes.”



#why Fuzzing ?

Fuzzing has a good **track record** of being efficient at bug finding

Compared to standard functional testing that focus on the “**happy path**”, fuzzing tries to cover the **edge cases**

Recommended by PCI as the de facto dynamic testing method



The bug-o-rama trophy case

Yeah, it finds bugs. I am focusing chiefly on development and have not been running the fuzzer at a scale, but here are uniquely interesting bugs that are attributable to AFL (in large part thanks to the work done by other users):

ICG ¹	libjpeg-turbo ^{1,2}	libpng ¹
libtiff ^{1,2,3,4,5}	mozjpeg ¹	PHP ^{1,2,3,4,5,6,7,8}
Mozilla Firefox ^{1,2,3,4}	Internet Explorer ^{1,2,3,4}	Apple Safari ¹
Adobe Flash / PCRE ^{1,2,3,4,5,6,7}	sqlite ^{1,2,3,4}	OpenSSL ^{1,2,3,4,5,6,7}
LibreOffice ^{1,2,3,4}	poppler ^{1,2}	freetype ^{1,2}
GnuTLS ¹	GnuPG ^{1,2,3,4}	OpenSSH ^{1,2,3,4,5}
PUTTY ^{1,2}	mpd ^{1,2}	nginx ^{1,2,3}
hash (post-shellshock) ^{1,2}	tcpdump ^{1,2,3,4,5,6,7,8,9}	JavaScriptCore ^{1,2,3,4}
pdfium ^{1,2}	ffmpeg ^{1,2,3,4,5}	libmatroska ¹
libarchive ^{1,2,3,4,5,6}	wirehark ^{1,2,3}	ImageMagick ^{1,2,3,4,5,6,7,8,9}
BIND ^{1,2,3}	QEMU ^{1,2}	lcms ¹
Oracle BerkeleyDB ^{1,2}	Android / libstagefright ^{1,2}	iOS / ImageIO ¹
FLAC audio library ^{1,2}	libstdc++ ^{1,2,3,4}	less / lesspipe ^{1,2,3}
strings (+ related tools) ^{1,2,3,4,5,6,7}	file ^{1,2,3,4}	dpkg ^{1,2}
rca ¹	systemd-resolved ^{1,2}	libyaml ¹
Info-Zip unzip ^{1,2}	libsan1 ^{1,2}	OpenBSD pftcl ¹

The Zen of Fuzzing



Fuzzing vs the world

Technique	Easy to start ?	Accuracy	Scalability
static analysis	easy	low	relatively good
dynamic analysis	hard	high	uncertain
symbolic execution	hard	high	bad
fuzzing	easy	high	good

Types of fuzzers

	Easy to start ?	Priori knowledge	Coverage	Ability to pass validation
Generation based	hard	needed, hard to acquire	high	strong
Mutation based	easy	not needed	low, affected by initial inputs	weak

What can a fuzzer find?

- Crashes, crashes, crashes !!!
- Bugs
 - Stack / Heap overflows
 - Null pointer dereferences
 - Off by ones
 - Uninitialized variables

The revolution: Coverage guided fuzzing



Leverage a feedback loop to guide input mutation

Coverage guided fuzzing principle

The Old Way: Stress test

```
while (1)
    input[0] = rand() % 20;
    input[1] = rand() % 20;

    if (input[0] % 2 == 0) {
        if (input[1] == 12) {
            printf("problem here\n");
        }
    }
    else {
        printf("no problem\n");
    }
}
```

The New Way: Feedback guided

```
while (1)
    input[0] = rand() % 20;
    input[1] = rand() % 20;

    if (input[0] % 2 == 0) {
        if (input[1] == 12) {
            printf("problem here\n");
        }
    }
    else {
        printf("no problem\n");
    }
}
```

FLAG1

FLAG2

FLAG3

Fuzzing process with AFL, the coverage guided ★ #1 fuzzer ★

C Sources

Compile with GCC

Binary

Run with input

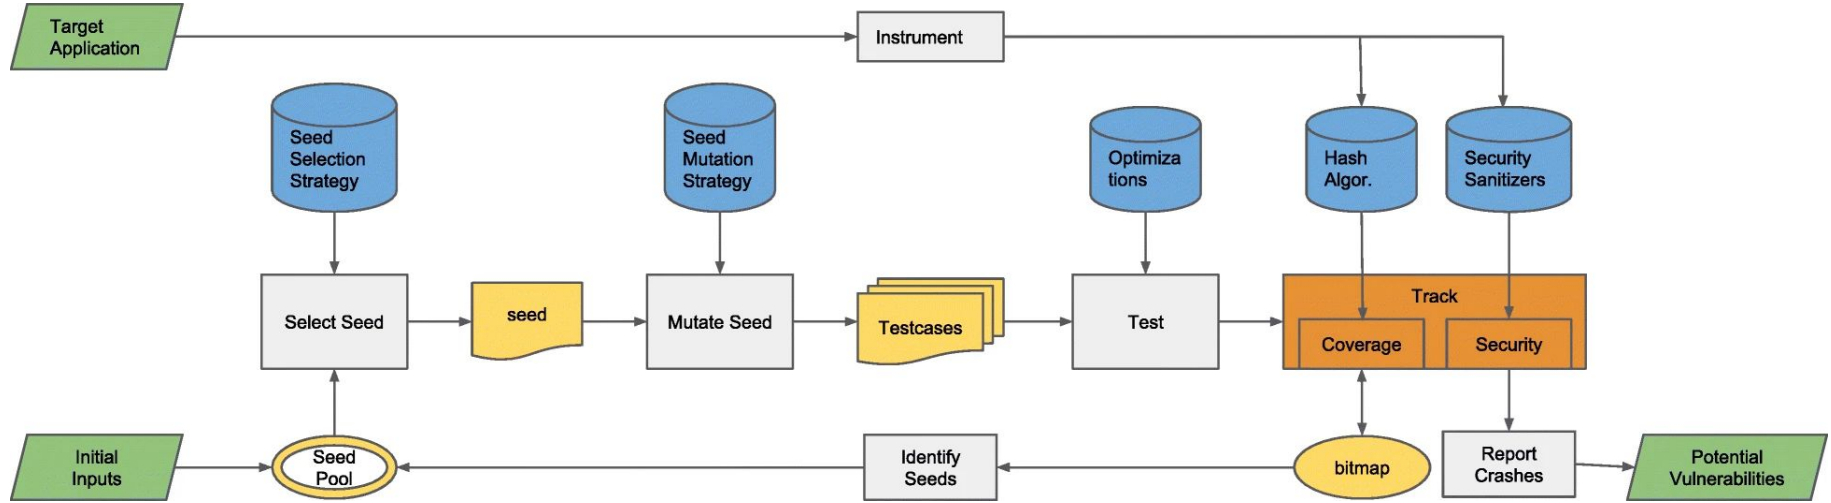
C Sources

Compile and *Instrument* with **AFL-GCC**

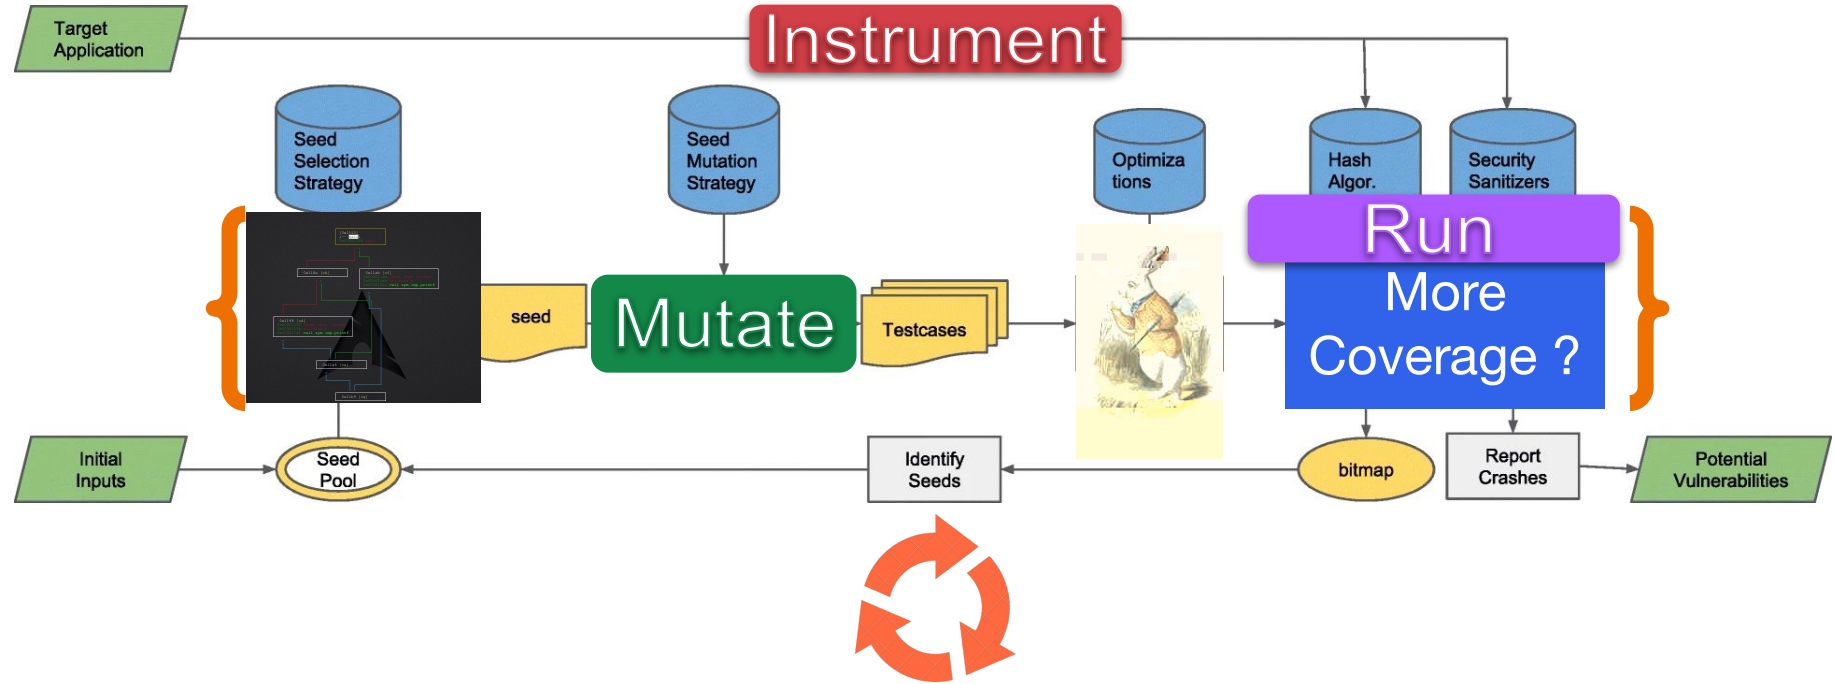
Instrumented Binary for Coverage

AFL-FUZZ Runs with *Mutated* input
guided by *Coverage*
& *Loop-Mutate until Maximize Coverage*


Fuzzer architecture



Fuzzer architecture explained



I know what you all think



Ahem ... <https://github.com/UnitTestBot/UTBotCpp>

Fuzzing exercise



Fuzzing exercise



Click here
Open VM
user/pass: fuzz/fuzz
Clone the repository in the VM
And get fuzzing !

First steps

```
# download VM (5.1 GB) https://drive.google.com/file/d/1UydxinlwKD847JHdbenO5qv7Xy7p3vJO/view
unzip Virtual\ Machine.zip

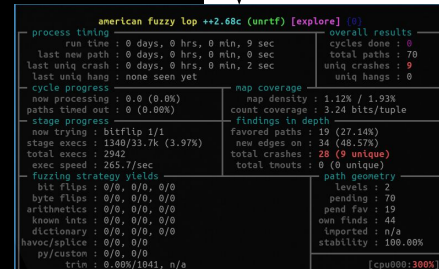
# install virtualbox on your OS
sudo apt-get install virtualbox

# or get it from here https://www.virtualbox.org/wiki/Downloads
# mac m1 users follow (next slide)

# run Virtualbox and file / import the image
# start the image and (recommended)
# take a snapshot after booting

# login with user/passwd : fuzz/fuzz and inside the image
cd ~/Desktop/WORKSHOP/Fuzz\ 0/unrtf
echo core | sudo tee /proc/sys/kernel/core_pattern
afl-fuzz -i ./tests -o afl-output -- ./bin/unrtf --verbose -P ./lib/unrtf/ @@
```

You should
get something
like this



```
american fuzzy lop ++2.68c (unrtf) [explore] 0.0
process timing:
  run time : 0 days, 0 hrs, 0 min, 9 sec
  last new path : 0 days, 0 hrs, 0 min, 0 sec
  last uniq crash : 0 days, 0 hrs, 0 min, 2 sec
  last uniq hang : none seen yet
cycle progress:
  now processing : 0.0 (0.00%)
  paths timed out : 0 (0.00%)
stage progress:
  now trying : bitflip 1/t
  stage execs : 1340/33.7k (3.97%)
  total execs : 2942
  exec speed : 265.7/sec
fuzzing strategy yields:
  bit flips : 0/0, 0/0, 0/0
  byte flips : 0/0, 0/0, 0/0
  arithmetics : 0/0, 0/0, 0/0
  known ints : 0/0, 0/0, 0/0
  dictionary : 0/0, 0/0, 0/0
  havoc/splice : 0/0, 0/0
  py/custom : 0/0, 0/0
  trim : 0.00%/1041, n/a
map coverage:
  map density : 1.12k / 1.93k
  count coverage : 3.24 bits/tuple
findings in depth:
  favored paths : 19 (27.14%)
  new edges on : 34 (48.57%)
  total crashes : 28 (9 unique)
  total tncouts : 0 (0 unique)
path geometry:
  levels : 2
  pending : 70
  pend fav : 19
  own finds : 44
  imported : n/a
  stability : 100.00%
[cpu000:300%]
```

Mac M1

install afl natively on mac m1

via rosetta 2

follow the two links below to install rosetta and afl x86 version on arm

<https://medium.com/mkdir-awesome/how-to-install-x86-64-homebrew-packages-on-apple-m1-macbook-54ba295230f>

<https://vineethbharadwaj.medium.com/how-to-compile-and-run-afl-fuzzer-on-m1-mac-with-apple-silicon-for-x86-instrumentation-support-4f1700eaafb6>

use rosetta to change arch

alias arm="env /usr/bin/arch -arm64 /bin/zsh --login"

alias intel="env /usr/bin/arch -x86_64 /bin/zsh --login"

for intel x86_64 brew

alias intelbrew='arch -x86_64 /usr/local/homebrew/bin/brew'

* intelbrew install afl-fuzz

* git clone https://github.com/antonio-morales/Hackfest_Advanced_Fuzzing_Workshop

* follow workshop slides

alternative method ... only as a last resort if previous solution did not work

* download and install UTM from https://mac.getutm.app/

* follow https://mac.getutm.app/gallery/ubuntu-20-04

* sudo apt install afl++

* git clone https://github.com/antonio-morales/Hackfest_Advanced_Fuzzing_Workshop

* follow workshop slides replacing afl-gcc with afl-clang-fast


First steps

```
cd ~/Desktop/WORKSHOP/Fuzz\ 0/unrtf
echo core | sudo tee /proc/sys/kernel/core_pattern
afl-fuzz -i ./tests -o afl-output -- ./bin/unrtf --verbose -P ./lib/unrtf/ @@

cd ~/Desktop/WORKSHOP/
git clone https://github.com/antonio-morales/Hackfest_Advanced_Fuzzing_Workshop
cd Fuzz1

# GOAL: Find the bug with AFL fuzzer
# Follow the hints
```

Follow the white rabbit



Embedded Fuzzing



```
# Run afl workshop inside docker emulating raspberry pi
```

```
docker run --network host -it -v $HOME/.dockerpi:/sdcard lukechilds/dockerpi  
// or use podman
```

```
# user / pass : pi / raspberry
```

```
# install afl / clone workshop and fuzz like previously
```

```
# GOAL: Find the bug with AFL fuzzer
```

```
# Follow the hints
```

